

A Conceptual Framework for Evaluating and Designing Information Discovery and Curation Web Tools

Leave Authors Anonymous
for Submission
City, Country
e-mail address

Leave Authors Anonymous
for Submission
City, Country
e-mail address

ABSTRACT

Everyday life involves the discovery and curation of digital information. People search the Web continuously, from quickly looking up information needed to complete a task, to endlessly searching for inspiration and knowledge. A variety of studies have modeled information seeking strategies and characterized curation activities on the Web. However, there is a lack of research on how existing Web applications support the discovery and curation of information, especially concerning user motivations and how different approaches can be compared. This paper presents a study of information discovery tools and how they relate to the nature of information seeking. We propose a conceptual framework of application design elements that support different aspects of information discovery and curation. This framework can be used for designing, evaluating and updating Web applications.

Author Keywords

Information discovery; information curation; Web design.

INTRODUCTION

Web technologies help people satisfy their information needs. People research their interests and hobbies using various online resources, shoppers search online stores for product characteristics to make purchasing decisions, and travelers visit online booking sites to find information about flights and hotels. To accommodate diverse and evolving user needs, Web applications continuously introduce new features and services, empowering information discovery and curation.

The term “information discovery” has been used to define or explain various information behaviour paradigms, such as information exploration [40] and serendipitous information seeking [14]. Information discovery can take on many forms. Web users might be hoping to find particular pieces of information, such as show times and phone numbers, to satisfy specific information needs [35]. Alternatively, they might be lacking well-articulated information needs, so they engage in opportunistic browsing [26]. Sometimes people discover information online without even looking for it [3]. The nature

of information discovery can vary, and therefore, requires elaborate tool support. The functionality required for information discovery and curation can also be distributed among multiple applications, which often leads to tools that provide integrated solutions. With people having such diverse information needs and methods of looking for information, designing for information discovery is a challenging task [9, 27].

Our research goal is to gain an understanding of how existing tools support digital information discovery and curation so that we can improve the design of Web applications for information discovery. While several researchers propose frameworks targeted at designing information discovery systems [35, 22], the importance of information curation in the realm of information discovery has been largely overlooked despite the rapidly increasing popularity of socially-curated information spaces. Moreover, much of the existing work that focuses on how people look for and discover information online [3, 6, 10, 20, 26, 29, 36] fails to examine concrete features of existing Web-based information discovery applications that empower real-world users. More research is necessary to determine how different tool features provide fundamental support for information discovery and curation.

To enhance information seeking and curating experiences and support users’ interactions, we extend existing research by: (1) deriving factors that enable information discovery and curation and relating them within a framework; (2) using the framework to establish a set of questions for evaluating and designing new applications; (3) iteratively evaluating the framework by using it to study and describe current Web applications, which in turn helped refine the framework of factors and questions; and (4) relating the framework to information discovery and curation motives that drive the underlying usage of Web-based applications.

WEB-BASED INFORMATION DISCOVERY AND CURATION

Given the complexity of Web-based information discovery and curation tasks, a variety of research topics are examined to gain an understanding of how current Web tools support these tasks, including information-related Web usage characteristics and current information behavior models.

Information Behavior

Information behavior refers to the totality of ways in which humans behave in relation to information [42]. A number of models and frameworks have attempted to represent human information behaviour in its entirety or to represent some of

Paste the appropriate copyright statement here. ACM now supports three different copyright statements:

- ACM copyright: ACM holds the copyright on the work. This is the historical approach.
- License: The author(s) retain copyright, but ACM receives an exclusive publication license.
- Open Access: The author(s) wish to pay for the work to be open access. The additional fee must be paid to ACM.

This text field is large enough to hold the appropriate release statement assuming it is single spaced.

Every submission will be assigned their own unique DOI string to be included here.

its components, such as information seeking and searching, information discovery, and information curation.

Information Seeking Models

Information seeking refers to “the purposive seeking for information as a consequence of a need to satisfy some goal [42].” Several researchers have tried to identify what different modes of information seeking behaviour may entail. According to Kellar et al. [20], information seeking is composed of browsing, fact finding, and information gathering. Although the authors categorized information gathering as part of information seeking, it appears to be more closely related to digital curation [5, 41]. Ellis et al. [10, 11, 12] proposed a model of information seeking characterized by six different patterns: starting, chaining, browsing, extracting, monitoring, and differentiating.

Information Foraging

Information foraging theory is another approach towards understanding how people adapt their strategies of interacting with technology when seeking, gathering, or consuming information, depending on the environment [34]. The theory resonates with explanations of human behavior in the context of food foraging. The underlying assumption of the information foraging theory is that people, similarly to when they forage for food, adopt their foraging strategies to the environment in order to gain the maximum amount of valuable information. The theory states that “natural information systems evolve towards stable states that maximize gains of valuable information per unit cost.” The theory introduces three key concepts to formulate an understanding of information foraging: 1) *information scent* that refers to proximal cues (often visual or linguistic) that people use to identify the value of information; 2) *information diet* which deals with user preferences when it comes to information and 3) *information patches* that are clusters of information that an information system presents before the user. This theory lays the foundation for existing information foraging models [16, 23] as well as social information foraging models [33, 15].

Information Discovery

Kerne and Smith proposed an information discovery framework [22] that connects human cognitive states to those of an information system. The framework represents a continuum of information flowing through different system and cognitive states as a result of an iterative reformulation process. The framework consists of five mental states: formulating a problem, evaluating results, updating and forming mental models, running mental models, and discovering solutions. Each mental state has a corresponding interaction with the system. For example, browsing resources (human-system interaction) facilitates evaluation of immediate results (cognitive state).

Digital Curation

In 2002, Bates extended her research on the topic of information behaviour with the notion of *information farming*, which involves people collecting and organizing information for future use and revisitation [4]. More commonly, information farming is referred to as digital curation. Whittaker suggests that in terms of Web use, a significant shift is happening from information consumption to information curation. People no

longer use the Web just to find and consume information of interest, but they also try to save and manage that information to refine and exploit later [41].

In summary, existing models and frameworks for information seeking, searching, exploration, discovery and curation try to explain human information-related behavior using different but comparable terminology. They help establish an understanding of how humans interact with information, however, they fail to address the required tool support for information-related activities or they address the needs at too high a level.

Web Tasks and Modes of Web Use

Outside the realm of cognitive models and frameworks for information behavior, we find research that examines information discovery, curation, and other Web information behaviours and corresponding tasks, methods, and modes.

Kellar et al. [20] separated Web tasks into five categories: transactions, browsing, fact finding, information gathering, and other uncategorized tasks. In later work, Kellar et al. [21] added communication and maintenance as additional Web tasks. Similarly, Sellen et al. [36] identified six tasks that are performed by Web users: browsing, finding, housekeeping, information gathering, communicating, and transacting. Using different terms, Kellar et al. and Sellen et al. both identified highly comparable tasks, such as fact finding and finding [information], housekeeping and maintenance, etc.

Building on Ellis’ model of information seeking [10, 11, 12], Choo et al. [6] derived anticipated Web tasks that correspond to the information seeking patterns in the model, such as identifying which Websites would point to information of interest, navigating through links, bookmarking, etc.

People often engage in information seeking activities to close some knowledge gap that occurred as a result of not having enough information to perform a task [35]. Therefore, when providing tool support for various information discovery tasks, it is useful to consider the motivations as they can be different for each task. Morrison et al. [29] make a distinction between methods of Web use and purposes. The authors derived a purpose-based taxonomy of Web use, including three purposes or motivations: finding information, comparing pieces of information or choosing products to make a decision, and using the Web to find relevant information to gain an understanding of some subject. Consequently, methods of finding information identified by Morrison et al. are collecting, finding, exploring, and monitoring. The differences between the two taxonomies suggest that different information seeking tasks may be performed to satisfy more than one information seeking purpose. Therefore, each purpose may require more than one task-supporting mechanism.

Categorizing Web usage into information seeking, digital curation, and other Web tasks may not adequately describe how information-related tasks are performed. Lindley et al. [26] conducted a qualitative study and identified five distinct modes of Web use: respite, orienting, opportunistic, purposeful, and lean-back. Understanding the characteristics of different modes can guide the design of Web interactions. For

example, opportunistic use can have unarticulated or continuously changing information needs. Later they may resume their opportunistic information seeking. Opportunistic use is also ‘grasshopper-like’ as users often jump from one resource to another [26]. Thus, there is a need to consider mechanisms for supporting users’ information needs, revisitation, and arbitrary navigation.

Different taxonomies of information seeking and curation tasks reflect on actual Web usage rather than theoretical modeling of human behavior, however, these taxonomies still focus on human activities when they interact with technology. A better understanding of how a system can support these activities is needed in order to effectively support human information-related interactions.

In summary, there are a multitude of tools that support different aspects of information discovery and curation, but understanding how these tools are similar (or differ) is difficult. Moreover, the existing research is not useful for identifying gaps in current tools or ways that current tools may be improved to support information discovery and curation. We address these problems by presenting a conceptual framework for information discovery and curation.

METHODOLOGY

Our methodology consisted of four steps. To gain a deeper understanding of the problem of information discovery and curation, we conducted an extensive literature review where we derived a preliminary set of information discovery and curation design factors and related them within a framework. The framework was then applied as part of the evaluation of 20 different information discovery applications and iteratively refined after every evaluation. Lastly, the framework was applied to a reevaluation of some of the previously evaluated tools with the purpose of validating its effectiveness.

Research Questions and Objective

This study was designed to address the problem of designing Web applications for information discovery and was motivated by the following:

RQ1: How do existing Web applications support information discovery?

RQ2: How do existing information discovery applications support information curation?

To address RQ1 and RQ2, we conducted an extensive literature review and a case study of 20 information discovery tools. Using insights from RQ1 and RQ2, we established our main research objective: **to develop a framework for performing summative and formative evaluation of Web-based information discovery and curation tools.**

Literature Review

The development of the framework began with an extensive literature review. A diverse set of topics contributed to forming an understanding of information discovery and curation, including information behaviour and information seeking models, high-level Web tasks, and modes of Web use,

exploration-based models of discovery, and methods of personal and social curation. From this review, the preliminary design factors for the framework were derived.

Building and Refining the Conceptual Framework

The Conceptual Framework we present in this paper was iteratively developed and refined through a careful analysis of 20 information discovery applications (see Table 1) and an in depth literature review. To guide the development of the framework, we selected some of the most used information discovery applications today. The popularity of information discovery applications was determined using Website popularity ranks provided by Alexa¹, a commercial Web traffic data provider. The focus was on applications that had strong information discovery components and less priority was given to applications whose purpose revolved only around curation. We examined the overall purpose of each application, its description as defined within the application, as well as literature and documentation related to the application (if they were available) against the features that the application provided. For example, if an application provided bookmarking features, we checked if it was indeed intended to be used for information preservation.

Our methodology involved an iterative process of selecting tools, analyzing them, and determining whether they could be described and evaluated using the framework. If we found a key feature that could not be described, we adapted the framework according to the findings. We repeated the process of tool selection and evaluation until the framework was usable for all tools. We then grouped the elements of the framework into categories, recording corresponding questions to ask in order to evaluate other applications.

Application	Address	Description
Pinterest	www.pinterest.com	Visual discovery tool
Delicious	delicious.com	Social bookmarking service
Tumblr	www.tumblr.com	Microblogging platform
StumbleUpon	www.stumbleupon.com	Web page discovery tool
Wikipedia	en.wikipedia.org	Free content Internet encyclopedia
Google Maps	www.google.ca/maps	Web mapping service
Rotten Tomatoes	www.rottentomatoes.com	Movie and TV database
500px	500px.com	Photography site
BucketList	bucketlist.org	Goal tracking and discovery service
We Heart It	weheartit.com	Visual discovery tool
Scoop.it!	www.scoop.it	Online publishing platform
Google Images	images.google.com	Image discovery service
Vimeo	vimeo.com	Video sharing Website
LifeHacker	lifehacker.com	Daily blog
YouTube	www.youtube.com	Video hosting platform
Yelp	www.yelp.ca	Business review site
IMDb	www.imdb.com	Movie database
Trip Adviser	www.tripadvisor.ca	Travel site
Urban Spoon	www.urbanspoon.com	Online bar and restaurant guide
Thesaurus	thesaurus.com	Online thesaurus

Table 1. Web-based Information Discovery and Curation Tools Studied

¹ Alexa is available at www.alexa.com

Framework Validation

In order to demonstrate the benefits of having such a framework, we applied the framework to five of the previously examined tools and show how its application can help in describing (and comparing) the features of these tools while revealing how the tools may be improved. Furthermore, we used the framework to guide the design of a novel tool. Due to space constraints, we can only report one of these results here (see Section Framework Validation) and we direct the reader to [38] for full details on this validation.

Limitations

The case study we conducted has a number of limitations. A lack of documentation, research literature, and formal descriptions of available features for some applications introduces a threat to the construct validity of the study. In addition, information discovery tools and features can be used in unintended or unforeseen ways by designers and developers. Therefore, the recorded use of some features within information discovery applications was recorded in our interpretations. To compensate for such limitations, the researchers personally employed the tools over an extended period of time to gain a deeper understanding of their use. In addition, we considered some cases with similar functionality and design to be able to validate or clarify prior findings. Additionally, framework validation was performed on five of the previously examined tools, introducing another limitation.

Another limitation was the lack of prior research on the subject matter. Some researchers have studied information seeking models and high-level Web tasks, but there is a lack of literature on how to enable and support different Web tasks. This opens up opportunities for future research to analyze methods of developing and building frameworks for facilitating and evaluating tools that support other Web tasks, such as communication, transactions, and goal realization.

A CONCEPTUAL FRAMEWORK FOR INFORMATION DISCOVERY AND CURATION ON THE WEB

In our framework, we build on existing models and frameworks of information discovery and curation and our analysis of existing Web tools to derive corresponding design factors for Web design. The first part of the framework deals with the *motives* behind information discovery and curation. These motives often define use cases for Web application design and help set initial assumptions about the required functionality.

The second part of the framework defines the *actions* that comprise discovery and curation activities, and the design factors that enable them. Some examples of actions include managing and preserving information. To support these actions, a Web-based application must provide corresponding mechanisms, such as bookmarking and tagging capabilities.

Actions can be further decomposed into *operations* performed using mechanisms that enable the actions. For example, the information preservation (action) can be enabled using a bookmarking feature (enabling mechanism) so that users can bookmark information using the feature (operation). The third part of the framework deals with improving operations for information discovery and curation.

Our framework considers human motives and relates information discovery and curation actions with corresponding enabling mechanisms. We relate operations that arise from actions with corresponding cognitive support, personalization, and automation. One of our very early versions of the framework [39], which has a different structure and an incomplete set of design features, also lacks the distinction between operations and actions, and the support they require. Similar terminology is used in Activity Theory [24] to describe human practices. Figure 1 gives a high-level overview of the framework and illustrates how the different components of the framework are connected.

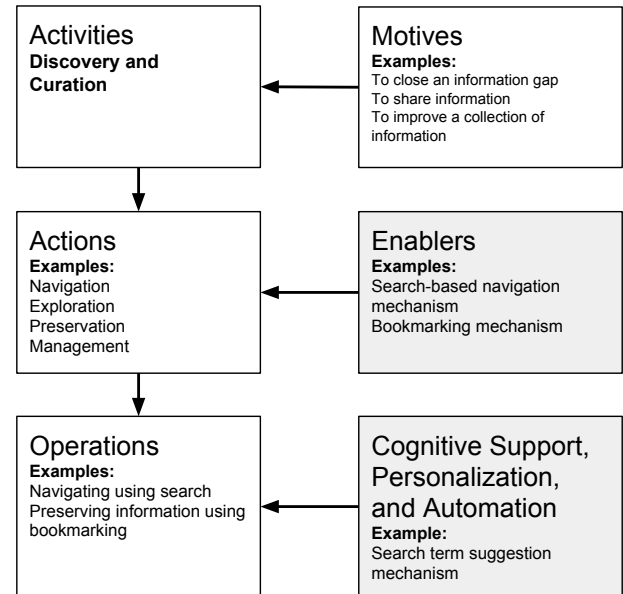


Figure 1. Framework Composition

Motives Behind Information Discovery and Curation

There are a wide variety of user motives behind information discovery and curation, and certain aspects of these motives can significantly impact the design of an application. Understanding a user's motives can help form a conceptual model of a needed Web application and its features. The following generalizations of motives and their properties can help define conceptual models and identify primary information discovery and curation use cases.

Closing a Knowledge Gap

The primary motive for information discovery is usually to close a knowledge gap that occurs when the user tries to accomplish a task but lacks information to do so. Depending on the context in which the motive arose, an information need can have various degrees of specificity. For example, if the motive is to find inspiration for a project, the information need is vaguely defined. However, if the motive is to find a phone number of a specific business, an information need is well-defined. In some cases, the information need may be hidden and the user might not be aware of the existing knowledge gap. The specificity of an information need determines important properties of information discovery mechanisms,

such as whether users can benefit more from mechanisms that allow them to specify an information need, help form an information need, or allow them to randomly retrieve information. This property has to be taken into consideration when evaluating or designing a Web application.

The nature of an information need predetermines whether discovery is serendipitous or oriented towards fact finding. Thus, an application can be designed to increase serendipity and opportunistic discovery or to improve purposeful fact finding. On the one hand, displaying featured content can improve serendipitous discovery because of its unexpected nature and novelty. On the other hand, using context (e.g., location and date) to tailor search results can improve fact finding.

Another motive type for information discovery relates to the two qualities of the Web defined by Lindley et al.[26]: *persistence* refers to the quality of the Web that allows people to habitually revisit Web pages and continue ongoing Web projects; and *temporality* refers to the quality of the Web that allows the content of Websites to be continuously updated to provide users with new information. Persistence alone usually facilitates information *rediscovery*, which is an act of refinding previously found information. However, if persistence is combined with temporality, they can facilitate discovery of new information within the same application or channel. We refer to this type of discovery as *channel-based discovery*. Some of the common motives for channel-based discovery include orienting (or monitoring for updates) and opportunistic information discovery [26].

The motive behind information rediscovery involves finding previously discovered information and reclosing the previously closed knowledge gap (e.g., in case the information was forgotten). It usually results in the user looking for previously found resources and Web pages. In fact, Web page revisitation is one of the most commonly performed Web browsing activities [2, 7]. The percentage of revisited Web pages involved in Web browsing can range from 58% [37] to 81% [8]. Some of the reasons for revisiting pages include shopping, communication, entertainment, education, activity planning, and hobby-related information retrieval [2]. Some Web pages and resources can be rediscovered using navigation, while others need to be previously preserved (bookmarked) to afford rediscovery. Rediscovery is one of the many ways in which information discovery and curation interweave.

Supporting Future Use and Reaccess

The main motive behind information curation is to make it possible to retrieve and use information. In order to facilitate easy information retrieval, many Web applications employ various forms of bookmarking systems. Traditionally, bookmarks must be manually organized into folders, but this method of organization is considered inefficient because folders with bookmarks become easily cluttered [1]. Therefore, in order to efficiently support information rediscovery, Web tools need to provide mechanisms for information preservation along with information management.

Improving Collections

People gather information to improve existing collections [26]. Although some deeper motives may include self reflection or the possibility of future use, collecting information is a motive in itself. Information gathering may be stretched over a period of time [20], resulting in repeated page visitation. Although information gathering comprises only 13.4% of Web usage, it contributes to many goal-supporting activities, such as decision making and planning [20].

Facilitating Communication

As part of his information behavior model, Wilson identified communication of information as an outcome of information seeking. Communication can also be thought of as a motive for information discovery and curation. To support communication of information, Web tools provide mechanisms that allow various users to share information among themselves. *Social bookmarking* is one popular way to preserve and share information across communities and to communicate with other users [13]. One of the first visions of social bookmarking was associated with Web blogging. Oravec [31] suggests that Web blogs help users annotate or bookmark important information and build a “map” of the Internet. The evolution of social bookmarking has led to advanced techniques for collaborative information discovery and curation.

In summary, while it is not feasible to list all of the possible motives for information discovery and curation, this section outlined some of the key motives that can aid in developing use cases and formalizing conceptual models for Web applications. These motives also make it easier to showcase how mechanisms for discovery and curation (presented in the next section) complement each other.

Discovery and Curation Activities

The next part of the framework deals with the actions associated with enablers of information discovery and curation. A more detailed overview is depicted in Figure 2; the two main activities (discovery and curation) are decomposed into actions, and each of the actions is supported by a group of features or mechanisms that enable a given aspect of discovery or curation in a Web application.

Navigation in Discovery: Following Information Scent

To discover information, a user needs a way to navigate to it. Navigation in information discovery can be thought of as following an information scent. In general, the information scent models deal with how users identify value, cost, or the access path of information sources based on proximal cues, such as links, icons, categories, etc. [34]. Common methods of navigation that facilitate information discovery include descriptive, referential, opportunistic, and system-regulated. We describe these methods below and direct the reader to the top part of Table 2 to see an overview of these methods and the enabling mechanisms possible for each method.

Descriptive Navigation: A navigation is descriptive when the user has a means of describing their information need. It is usually implemented as *search-based navigation* since it allows users to enter a search query and describe what they want to find. Some modern descriptive navigation systems are voice-activated.

Descriptive navigation can also help to rediscover information, but it is not always a reliable way of refinding information [7]. In information portals that provide access to fairly ambiguous information and that have regularly updated information flow, the search feature is usually designed around retrieving information related to some general topic. In order to make search-based navigation a reliable way to rediscover information, it must return consistent results.

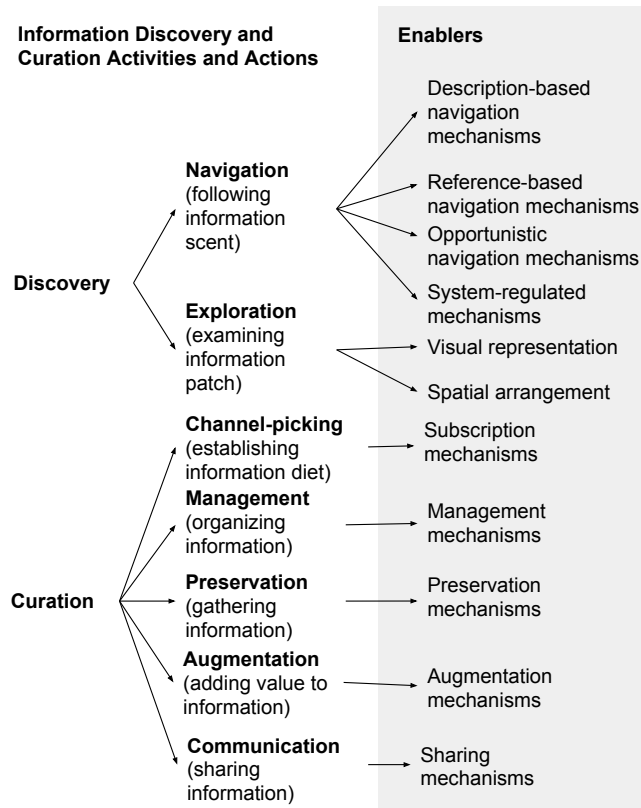


Figure 2. Information Discovery and Curation Activities, Actions, and Corresponding Enablers

Referential Navigation: A navigation is referential when the user finds a reference to the term that they are looking for, such as a link or icon. This reference represents an information scent. The underlying assumption of this method of navigation is that the user can recognize the needed information or a reference to it as they see it [40].

Referential navigation mechanisms can take many forms. Some common types are categories, facets, filters and tags. In some applications, users can search by a given resource. For example, YouTube provides a playlist with music related to the currently playing song. Information scent representatives may also reference sources outside of the given system, enabling another type of integration of Web applications. Referential navigation can help the user identify their information needs by suggesting terms, topics or categories to use, and therefore, direct the user to relevant resources [25]. It can also help narrow the results to a specific type of resource so that further discovery is bounded by that type. For example, TripAdvisor helps

narrow search results by allowing users to choose among hotels, flights, vacation rentals, restaurants and destinations.

Opportunistic Navigation: Opportunistic navigation is a method of navigating ‘randomly’ through resources and Web pages. We call this ‘opportunistic’ because it is not truly random, but its serendipitous nature makes users feel like it is. This navigation method is especially useful when the information need is fully undefined. Many applications support opportunistic jumping from one resource to another. For example, StumbleUpon makes it possible to explore the Web in general—other Websites and Web applications, allows for integrated navigation—whereas Wikipedia provides opportunistic access to its own articles.

System-regulated Navigation: Web applications often display information without the user’s active participation. This information can be a news feed, featured deals or articles or other types of content. We refer to this type of navigation as system-regulated because it occurs when the application brings the content to the user instead of the user applying any effort to find content. It differs from opportunistic navigation because the the user cannot choose when to observe new information; instead, all updates are regulated by the application. One application that utilizes system-regulated navigation is Yelp. As soon as the user enters the site, this tool displays featured restaurants as well as the user’s recent activities. As with any other navigation method, system-regulated navigation can ensure cross-application integration by displaying content from other Web applications.

Exploration in Discovery: Examining Information Patches

Exploration of resources is another action that facilitates information discovery. Visual and spatial cues, which help representing single or multiple resources, enable this action by allowing users to conveniently examine information patches (please refer to the middle part of Table 2).

Visual and textual previews: Abrams et al. [1] identified link representation as one of the problems with traditional bookmarking. Analogous with browsing through a bookmark manager, identifying relevant information when browsing through links in a Web application can be a challenging task. Visual previews and textual previews make it easier to evaluate the relevance of resources by providing the user with more information scent. Many social bookmarking systems, such as Scoop.it! and Pinterest, support visual previews of bookmarked pages. Delicious is a social bookmarking application that lacks this type of link representation support, and so it is harder to determine if a link will lead to a relevant resource.

Visual and textual cues: Visual and textual information cues are also important for a single resource exploration. Not only do they help navigation within the resource or Web page, but they can also contribute to the learning experience. For example, if the user would like to know what something looks like, they can learn it from the representation in question.

Spatial visualizations: Similar to link representation, spatial visualization of numerous links is another problem that oc-

Types of discovery and curation features	Questions to be posed during the design or evaluation of discovery and curation tools and sample features	
Navigation features	Enabling mechanisms	Cognitive support, personalization, and automation
Descriptive	How does the application support descriptive navigation? Search-based navigation / Integrated search	How can descriptive navigation mechanisms be enhanced? Personalized results / Guided search
Referential	How does the application support referential navigation? Categories / Facets / Filters / Tags / Search by item or resource / Integrated reference	How can referential navigation mechanisms be enhanced? Suggesting categories / Suggesting topics of interest / Suggesting tags / Suggesting similar resources
Opportunistic	How does the application support opportunistic navigation? Opportunistic navigation feature / Integrated opportunistic navigation	How can opportunistic navigation mechanisms be enhanced? Personalized opportunistic navigation
System-regulated	How does the application support system-regulated navigation? Static direct display / Integrated static display / Featured content / Integrated featured content / News feed / Integrated news feed	How can system-regulated navigation mechanisms be enhanced? Personalized featured content / User activity update notification / Application activity update notification / Artifact update notification
Exploration features	Enabling mechanisms	Cognitive support, personalization, and automation
Visual and textual cues of multiple resources	How does the application use visual and textual cues to help identify resources of value? Visual preview / Textual preview	How can visual and textual cues of multiple resources be enhanced? Personalized visual preview / Personalized textual preview
Visual and textual cues of a single resource	How does the application use visual and textual cues to help identify the value of information within a resource? Visual cues / Textual cues	How can visual and textual cues of a single resource be enhanced? Personalized visual cues / Personalized textual cues
Spatial proximal cues of multiple resources	How does the application use spatial proximal cues to effectively present multiple resources? List / Grid / Gallery / Spatial semantic / Consistency	How can the use of spatial proximal cues be enhanced to effectively present multiple resources? Personalized arrangement of multiple resources
Spatial proximal cues of a single resource	How does the application use spatial proximal cues to effectively present information within a single resource? Spatial semantic / Consistency	How can the use of spatial proximal cues be enhanced to effectively present information within a resource? Personalized arrangement of information within a resource
Curation features	Enabling mechanisms	Cognitive support, personalization, and automation
Management	How does the application support management of information? Public or private collection-based categorization / Public or private tag-based categorization	How can management of information be enhanced? Suggesting collections / Suggesting tags / Automated classification into collections / Automated tagging
Preservation	How does the application support preservation of information? Internal preservation of internal resources / Internal preservation of external resources / External preservation of internal resources	How can preservation of information be enhanced? History / Suggested preservation
Augmentation	How does the application support augmentation of information? Annotation / Evaluation	How can augmentation of information be enhanced? Automated augmentation / Suggested augmentation
Sharing	How does the application support sharing of information? Adding resources / Internal sharing / External sharing	How can sharing of information be enhanced? Automated sharing / Suggested sharing
Channel-picking	How does the application support the user in establishing their information preferences? User subscription / Site subscription / Artifact subscription	How can establishment of information preferences be enhanced? Suggesting users for subscription / Suggesting artifacts for subscription / Automated subscription

Table 2. Types of Information Discovery and Curation Features and Related Questions

curs when browsing through diverse content [1]. Therefore, a semantic to the spatial arrangement of information (single and multiple resources) is of major importance. Information discovery applications often employ sophisticated ways of spatially arranging resources to make it easier to browse through large amounts of information. Common ways of arranging multiple resources include `list`, `grid`, and `gallery` layouts. Additionally, consistency in the way multiple and single resources are represented helps form a conceptual model of how the application can be used and provides some degree of predictability [30].

Curation

Information curation is a common activity across many information discovery applications. By asking questions about application design with regards to information curation de-

signers can find ways to add value to information and enable information discovery over time. Information discovery applications vary from being completely socially curated and populated by users, to those that lack any curation mechanisms. By definition, digital information curation is the notion of managing, preserving, and adding value to collections of information [5, 41]. Thus, curation activity consists of actions such as information management, preservation, information augmentation, sharing, and channel-picking. Refer to the bottom of Table 2 for this part of the framework.

Information management is one of the key elements of information curation [5, 41]. These mechanisms are prevalent in applications that have a lot of information that is hard to categorize automatically or can mean something different for each user. In the context of Web information management, `tag` and `collection`-based information categorization

play major roles. Resource categorization also helps establish relationships between various resources [5, 41]. Tagging can aid rediscovery and discovery in a socially curated space, as well as add more value to resources [18]. Sample applications that facilitate information management are Pinterest, a tool that supports tagging and collection-based categorization, and Tumblr, a tool that supports tagging.

Information preservation is a common Web task usually performed with the intent of revisiting information [1, 41]. However, information gathering is sometimes performed with just the goal of collecting information [26]. Bookmarking is a traditional way of preserving information. Internal preservation of internal resources means bookmarking resources can be reaccessed within the same application. Such bookmarking facilitates information curation within the system. Internal preservation of external resources signifies bookmarking other Web pages within an application. External preservation means bookmarking resources so that they are available through other bookmarking systems. An application must facilitate integration with other applications to enable external preservation [1].

Augmentation: One of the most important elements of digital curation is augmentation: adding value to information [5, 41]. It is often performed within social bookmarking systems, and many Web applications allow users to add value to the resources they curate. One way to augment information is by annotating it with comments and descriptions. Annotations are metadata attached to a resource that make it easier to search for and interpret information. For example, Yelp and TripAdvisor largely rely on reviews written by their users. Evaluation methods can have various forms. They usually take place in socially curated information systems. However, evaluation can also contribute to personal reflection and information preservation. Many applications allow users to evaluate resources by rating them or recording other forms of approval or disapproval, such as "I like this" and "I dislike this" buttons on YouTube.

Sharing information is key to empowering social information curation [5]. Therefore, the main components that facilitate sharing are the adding of resources, and external and internal information sharing. Adding resources not only facilitates global Web information curation, but it also scales the information available through the system, providing more opportunities for information discovery. Resources can be created by users themselves, taken from some other sources online, or both. For example, YouTube allows users to upload their own videos, whereas Pinterest permits adding images from other sites in addition to users' personal images. Sharing resources through different media and resharing them within the Web application supports channel-based information discovery within the media channels. Information discovery applications commonly allow for sharing information on popular networking sites outside the application.

Channel-picking is an action of selecting information sources. A common enabler for this action is subscrip-

tions that help users follow the news [19]. To support channel-based discovery, an application must provide a subscription mechanism. For example, Rotten Tomatoes allows subscriptions to newsletters, but it does not allow subscriptions to movie critics that would be allowed with a user-based subscription mechanism such as the one in Pinterest. In some applications, the content is updated and curated by users, and users can subscribe to other users or artifacts. Similar to site subscriptions, user and artifact subscriptions are subscriptions to activity updates. These subscription mechanisms help with networking and provide awareness about other users' activities [28]. Such subscriptions also help filter new content delivered to the user.

In summary, information discovery and curation tools can have different implementations depending on the motives behind the activities. The design factors presented in this section can help enable different actions associated with information discovery and curation. However, the activities can be significantly improved by additional support and automation, as described in the next section.

Enhancing the Discovery and Curation Experience

The information discovery and curation enablers just presented are design elements that afford various operations. For example, the search feature enables typing in a query and searching for information. These operations can be further supported by another set of design elements that introduce cognitive support, personalization, and automation. The primary goal of this part of the framework (see right column of Table 2) is to highlight opportunities for improvement over various information discovery and curation enabling mechanisms. Strategies for improvement include providing additional cognitive support for a given operation, personalizing the user experience, and automating an operation. Not all of the strategies are feasible for every single operation, and some operations can be supported in multiple ways. The following outline some of the possibilities for advancing information discovery and curation features.

Enhancing Navigation: There are two common methods of enhancing information discovery when search-based navigation is used (see Table 2). The first method entails returning personalized results when the user enters a search query. Personalization can be accomplished using a variety of techniques, including predefined user preferences, social interactions, context, browsing history, etc. The second method is to suggest search terms to make it easier for the user to formulate their information need. For example, Yelp suggests search terms as the user enters their query.

To further support referential navigation, applications can personalize reference suggestions, such as categories, tags, and topics of interest. They can also suggest relevant resources based on the one that the user already selected. As an example, after a user clicks a 'pin', Pinterest showcases other similar 'pins'.

For opportunistic navigation, Web tools sometimes allow users to personalize types or categories of information that they the users would like to discover. StumbleUpon al-

lows users to not only choose topics of interests, but it can also help them discover new promising topics.

Featured content can also be personalized to improve information discovery with system-regulated navigation. For example, Yelp showcases restaurants from a predefined area, such as the city where the user is from.

Finally, to make better use of subscribed content and to reduce human efforts when searching, an application can support various notification mechanisms. These can advise the user about updates on the Website content, various artifacts, and activities of other users.

Enhancing Exploration: Personalization of the spatial information representation usually has limited support in Web applications. Presumably, it is because consistency is more welcomed within information discovery applications than spatial personalization. However, it is still possible to personalize the arrangement of multiple resources or information within a single resource. Visual and textual personalizations are more common, especially when the content within the application is curated by its users. For example, Flickr Web application for managing and sharing photographs personalizes album covers so that they are easier to rediscover.

Enhancing Curation: Information management can be improved if the system helps the user make decisions about information categorization or tagging (see Table 2). Alternatively, information can be categorized or tagged automatically. For example, when the user bookmarks a restaurant on Yelp, it is automatically categorized. The user can filter bookmarks by category whenever they go into the embedded bookmark manager.

Preservation operations can also be automated. An example of the most common automatic preservation mechanism is history. Applications such as YouTube and Google Maps preserve users' browsing history so that they can review it later. Additionally, preservation mechanisms can be suggested to the user. YouTube allows users to automatically share information about their activities, such as comments, added videos, liked or disliked videos, and created playlists. In general, socially curated spaces offer sharing channels to support convenient information communication. Augmentation is another aspect of information curation that can be either automated for or suggested to the user. For example, Yelp asks users to rate the places which the application identifies as having been visited by the user.

Notification mechanisms enable user awareness about new content on the subscribed channel [28]. Web applications that facilitate rapidly updating content support various notification mechanisms, such as messages within the application, informative emails, and smartphone notifications. Many types of notifications include suggested users or artifacts to follow. Some Web tools automatically subscribe users to notifications, usually during the registration process.

In summary, providing cognitive support, personalization, and automation dramatically improves the user experience when people interact with information discovery and curation systems. The framework can be used for identifying gaps in information discovery support and developing new technologies as described in the following section.

FRAMEWORK VALIDATION

In order to validate the conceptual framework and verify its stability, we used it to design a tool for location based photograph discovery and curation, we used it to evaluate five of the applications that were used in the construction of the preliminary framework: Pinterest, Google Maps, Wikipedia, Delicious, and Yelp. Due to space limitations, we only present the evaluation for one of the tools, Pinterest. The other tool evaluations and the location based discovery tool description can be found in [38]. As for Pinterest, we first summarize our observations resulting from asking the questions from the framework in a systematic manner. Based on our assumptions, judgment, and use of the framework, we propose directions for future development and reflect on certain needed mechanisms, as not all mechanisms are always required.

Pinterest is a Web application designed for image discovery and curation, oriented towards finding inspiration and collecting knowledge about hobbies and interests [17, 43, 32]. Users of Pinterest are commonly referred to as 'pinners'. Resources on Pinterest are called 'pins', and each 'pin' consists of an image, a short description, the user's name, and the name of the collection that the pin belongs to. More information is available once the user clicks on a 'pin'.

Motivated by the desire to gain inspiration and knowledge, Pinterest users have either under-defined or absent information needs. Other motives for using Pinterest could be to rediscover previously found information (and possibly use it), to be oriented about new 'pins' that emerge from subscribed channels, and to gather information for future rediscovery and the act of collection itself.

Navigation in Pinterest is mostly supported by descriptive, referential, or system-regulated mechanisms. Although an explicit opportunistic navigation mechanism is absent, both descriptive and referential mechanisms usually return novel and serendipitous results to facilitate opportunistic browsing. Descriptive navigation is enabled with a guided search mechanism that suggests search terms to the user.

Referential navigation is enabled in Pinterest using a range of techniques. To support articulation of an information need, a category-based navigation mechanism makes suggestions on subcategories or interests. Through clicking on a 'pin', the user can see related resources, enabling resource-based referential navigation. Most of the images on Pinterest are 'pinned' from other Websites, and users are provided with links to their original sources. Therefore, Pinterest supports integrated referential navigation.

System-regulated navigation within Pinterest is highly personalized. When the user enters the site, they see a history of their own information gathering activities and updates from

the people they are subscribed to. Additionally, the application suggests featured ‘pins’ based on the user’s interests.

To reinforce the discovery of visual data, Pinterest provides extensive support for various exploration mechanisms. Multiple resources are represented in a gallery layout, often referred to as a ‘pinboard’. This layout provides good spatial support for exploration and makes it easier to build a mental model of the tool by drawing analogies with a real pinboard. Users can create multiple ‘pinboards’ (also known as ‘boards’) which have personalized covers to enhance future exploration and rediscovery. A single resource does not have a lot of distinct spatial arrangements, however, it provides a visual glimpse into what can be found on the Website that the image came from, with textual preview being limited to the address of the Website.

Information management is accomplished through sorting ‘pins’ into different collections (‘pinboards’) thus enabling collection-based classification and internal preservation of internal and external resources. All user information collecting actions are automatically preserved and displayed. Users can augment the information pool by uploading new ‘pins’, commenting on existing ‘pins’, or adding descriptions. Users can also internally share ‘pins’ among themselves. Channel-picking actions are carried out by following or subscribing to users or individual ‘pinboards’. The system also automatically sends notifications through emails and suggests new ‘boards’ to follow.

Applying the framework to Pinterest revealed that it employs a variety of techniques to facilitate information discovery and curation. However, individual mechanisms could be further improved. For example, textual previews of multiple and individual resources is rather limited and provides little insight into what information source Websites actually contain. Furthermore, Pinterest could benefit from automatically classifying ‘pins’ into ‘boards’ because finding an appropriate ‘board’ for a ‘pin’ can be difficult when a user has a large number of existing ‘boards’. Overall, Pinterest provides rich support for information discovery and curation, and in some ways, enables each of the discovery or curation actions of the conceptual framework.

RESEARCH AND DESIGN IMPLICATIONS

The conceptual framework for information discovery and curation is designed to perform formative and summative evaluation of existing Web applications and to reveal how these tools support information-related activities in question. The framework as a tool and its ability to guide the process of analyzing Web applications makes it broadly applicable in research and Web design.

In the previous section, we demonstrated how the framework can be used to reveal missing features in tools. Using similar methods, the framework can also be applied to compare different Web applications. When used for evaluation, the framework helps to identify which areas of a tool require further attention. Therefore, the framework can be helpful for

designers who wish to improve existing tools or get ideas for new information discovery and curation applications.

Factors and questions of the framework are there to guide the developer and may expose gaps, but they do not dictate which features should be in an application. It is up to designers to decide whether to close those gaps and some gaps cannot be closed because of certain constraints or trade-offs that have to be made, such as data type and system design.

Even though applying the framework requires initial expertise and critical reasoning, it opens up opportunities for research and practice. For the research domain, the framework can serve as a guide for drawing distinctions between different Web-based information discovery and curation applications, finding gaps in tools, and selecting cases for studies based on required functionality. While, systematic evaluation of Web tools for information discovery and curation helps the designer improve user experience and gain better understanding of information behaviour within a given system.

FUTURE WORK AND CONCLUSIONS

In our study, we analyzed information curation and seeking tasks and developed a conceptual framework of factors and questions that are important when building and evaluating Web information discovery and curation tools. We then evaluated and iteratively refined the framework by analyzing 20 different information discovery applications and provided concrete examples of tool support addressing various concepts of the framework. Finally, we validated the framework by reevaluating five of the previously examined tools and used it to design a novel application (described in [38]).

The current version of the framework is designed to be generally applicable to information discovery applications. Finding ways to instantiate the framework and extend it for use in domain-specific practices could serve as a potential future research goal. For example, video discovery and curation activities have unique properties related to the type of data to be discovered—information is mostly found in the video itself, and it cannot be viewed all at the same time. Hence, the framework could be extended to address domain-specific challenges. Another potential research direction would be to expand our investigation to include factors that influence the need for one information discovery type over another and further deepen an understanding of the relationships between the motives for information discovery and curation activities and information discovery types.

Our framework opens up opportunities for structured information discovery and curation tool evaluation and design. As more tools are being developed within the social space of information discovery and curation, understanding how these tasks can be supported promises advancements in how Web applications are designed.

REFERENCES

1. David Abrams, Ron Baecker, and Mark Chignell. 1998. Information archiving with bookmarks: personal Web space construction and organization. In *Proceedings of the SIGCHI conference on Human factors in computing systems*. ACM Press/Addison-Wesley Publishing Co., 41–48.
2. Eytan Adar, Jaime Teevan, and Susan T Dumais. 2008. Large scale analysis of web revisitation patterns. In *Proceedings of the SIGCHI conference on Human Factors in Computing Systems*. ACM, 1197–1206.
3. Marcia J Bates. 1986. An exploratory paradigm for online information retrieval. *Intelligent Information Systems for the Information Society*. Amsterdam: North-Holland (1986), 91–99.
4. Marcia J Bates. 2002. Toward an integrated model of information seeking and searching. *The New Review of Information Behaviour Research* 3 (2002), 1–15.
5. Neil Beagrie. 2008. Digital curation for science, digital libraries, and individuals. *International Journal of Digital Curation* 1, 1 (2008), 3–16.
6. Chun Wei Choo, Brian Detlor, and Dan Turnbull. 2000. Information seeking on the Web: An integrated model of browsing and searching. *first monday* 5, 2 (2000).
7. Andy Cockburn, Saul Greenberg, Steve Jones, Bruce McKenzie, and Michael Moyle. 2003. Improving Web page revisitation: Analysis, design, and evaluation. (2003).
8. Andy Cockburn and Bruce McKenzie. 2001. What do Web users do? An empirical analysis of Web use. *International Journal of human-computer studies* 54, 6 (2001), 903–922.
9. AM Conaway, CK Pikas, UE McLean, SD Morris, LA Palmer, L Rosman, SA Sears, E Uzelac, and SM Woodson. 2010. Designing for information discovery: User needs analysis. *Johns Hopkins Applied Technical Digest* 28, 3 (2010), 290–291.
10. David Ellis. 1989. A behavioural model for information retrieval system design. *Journal of information science* 15, 4-5 (1989), 237–247.
11. David Ellis, Deborah Cox, and Katherine Hall. 1993. A comparison of the information seeking patterns of researchers in the physical and social sciences. *Journal of documentation* 49, 4 (1993), 356–369.
12. David Ellis and Merete Haugan. 1997. Modelling the information seeking patterns of engineers and research scientists in an industrial environment. *Journal of documentation* 53, 4 (1997), 384–403.
13. Enrique Estellés, Esther Del Moral, and Fernando González. 2010. Social bookmarking tools as facilitators of learning and research collaborative processes: The Diigo case. *Interdisciplinary Journal of E-Learning and Learning Objects* 6, 1 (2010), 175–191.
14. Allen Foster and Nigel Ford. 2003. Serendipity and information seeking: an empirical study. *Journal of Documentation* 59, 3 (2003), 321–340.
15. Wai-Tat Fu. 2008. The microstructures of social tagging: a rational model. In *Proceedings of the 2008 ACM conference on Computer supported cooperative work*. ACM, 229–238.
16. Wai-Tat Fu and Peter Pirolli. 2007. SNIF-ACT: A cognitive model of user navigation on the World Wide Web. *Human-Computer Interaction* 22, 4 (2007), 355–412.
17. Eric Gilbert, Saeideh Bakhshi, Shuo Chang, and Loren Terveen. 2013. I need to try this?: a statistical overview of pinterest. In *Proceedings of the SIGCHI conference on human factors in computing systems*. ACM, 2427–2436.
18. Thomas Gruber. 2007. Ontology of folksonomy: A mash-up of apples and oranges. *International Journal on Semantic Web and Information Systems (IJSWIS)* 3, 1 (2007), 1–11.
19. Akshay Java, Pranam Kolari, Tim Finin, Anupam Joshi, and Tim Oates. 2007. Feeds That Matter: A Study of Bloglines Subscriptions.. In *ICWSM*.
20. Melanie Kellar, Carolyn Watters, and Michael Shepherd. 2006. A Goal-based Classification of Web Information Tasks. *Proceedings of the American Society for Information Science and Technology* 43, 1 (2006), 1–22.
21. Melanie Kellar, Carolyn Watters, and Michael Shepherd. 2007. A field study characterizing Web-based information-seeking tasks. *Journal of the American Society for Information Science and Technology* 58, 7 (2007), 999–1018.
22. Andruid Kerne and Steven M Smith. 2004. The information discovery framework. In *Proceedings of the 5th conference on Designing interactive systems: processes, practices, methods, and techniques*. ACM, 357–360.
23. Muneo Kitajima, Marilyn H Blackmon, and Peter G Polson. 2000. A comprehension-based model of Web navigation and its application to Web usability analysis. In *People and Computers XIV Usability or Else!* Springer, 357–373.
24. Kari Kuutti. 1996. Activity theory as a potential framework for human-computer interaction research. *Context and consciousness: Activity theory and human-computer interaction* (1996), 17–44.
25. Mark Levene. 2011. *An introduction to search engines and web navigation*. John Wiley & Sons.
26. Siân E Lindley, Sam Meek, Abigail Sellen, and Richard Harper. 2012. It's simply integral to what I do: enquiries into how the web is weaved into everyday life. In *Proceedings of the 21st international conference on World Wide Web*. ACM, 1067–1076.

27. Gary Marchionini. 2006. Exploratory search: from finding to understanding. *Commun. ACM* 49, 4 (2006), 41–46.
28. David Millen, Jonathan Feinberg, and Bernard Kerr. 2005. Social bookmarking in the enterprise. *Queue* 3, 9 (2005), 28–35.
29. Julie B Morrison, Peter Pirolli, and Stuart K Card. 2001. A taxonomic analysis of what World Wide Web activities significantly impact people’s decisions and actions. In *CHI’01 extended abstracts on Human factors in computing systems*. ACM, 163–164.
30. Donald A Norman. 2002. *The design of everyday things*. Basic books.
31. Jo Ann Oravec. 2002. Bookmarking the world: Weblog applications in education. *Journal of Adolescent & Adult Literacy* (2002), 616–621.
32. Raphael Ottoni, Joao Paulo Pesce, Diego B Las Casas, Geraldo Franciscani Jr, Wagner Meira Jr, Ponnuram Kumaraguru, and Virgilio Almeida. 2013. Ladies First: Analyzing Gender Roles and Behaviors in Pinterest.. In *ICWSM*.
33. Peter Pirolli. 2009. An elementary social information foraging model. In *Proceedings of the SIGCHI Conference on Human Factors in Computing Systems*. ACM, 605–614.
34. Peter Pirolli and Stuart Card. 1999. Information foraging. *Psychological review* 106, 4 (1999), 643.
35. Henderik Alex Proper and PD Bruza. 1999. What is information discovery about? *Journal of the American Society for Information Science* 50, 9 (1999), 737–750.
36. Abigail J Sellen, Rachel Murphy, and Kate L Shaw. 2002. How knowledge workers use the web. In *Proceedings of the SIGCHI conference on Human factors in computing systems*. ACM, 227–234.
37. Linda Tauscher and Saul Greenberg. 1997. How people revisit web pages: Empirical findings and implications for the design of history systems. *International Journal of Human-Computer Studies* 47, 1 (1997), 97–137.
38. Elena Voyloshnikova. 2015. *A Conceptual Framework for Evaluating and Designing Information Discovery and Curation Tools*. Master’s thesis. University of Victoria.
39. Elena Voyloshnikova and Margaret-Anne Storey. 2014. Towards Understanding Digital Information Discovery and Curation. In *Proceedings of the 2014 Conference of the Center for Advanced Studies on Collaborative Research*. ACM, 247–261.
40. John A. Waterworth and Mark H. Chignell. 1991. A model of information exploration. *Hypermedia* 3, 1 (1991), 35–58.
41. Steve Whittaker. 2011. Personal information management: from information consumption to curation. *Annual review of information science and technology* 45, 1 (2011), 1–62.
42. Thomas D Wilson. 2000. Human information behavior. *Informing science* 3, 2 (2000), 49–56.
43. Michael Zarro and Catherine Hall. 2012. Pinterest: Social collecting for# linking# using# sharing. In *Proceedings of the 12th ACM/IEEE-CS joint conference on Digital Libraries*. ACM, 417–418.